

Homework #3: SQL

To start with, if you felt the class was unclear, check out the following tutorial: <https://mode.com/sql-tutorial/introduction-to-sql/>

Now! We'll be using sqlite to access a database. Start by downloading the sql lite file and putting it in the same directory as this notebook: <https://www.kaggle.com/datasets/kaggle/sf-salaries> (hit the 'download' button in the upper right). Check out the description of the data so you know the table / column names.

The following code will use sqlite to create a database connection.

```
import sqlite3
import pandas as pd

conn = sqlite3.connect("database.sqlite")
crsr = conn.cursor()
```

Exploration

Problem 1:

Try to create a query that gives you a data frame of the EmployeeName, JobTitle, and BasePay from the salaries table.

```
query = 'SELECT EmployeeName,JobTitle,BasePay FROM salaries'

df = pd.read_sql(query, conn)
df.head()
```

	EmployeeName o...	JobTitle object	BasePay object	
0	NATHANIEL FORD	GENERAL MANAGER-...	167411.18	
1	GARY JIMENEZ	CAPTAIN III (POLICE...	155966.02	
2	ALBERT PARDINI	CAPTAIN III (POLICE...	212739.13	
3	CHRISTOPHER CHONG	WIRE ROPE CABLE MAINTENANCE...	77916	
4	PATRICK GARDNER	DEPUTY CHIEF OF DEPARTMENT,...	134401.6	

Problem 2.

Modify your query to limit it to the year 2012.

```
query = 'SELECT * FROM salaries WHERE Year = 2012'

df = pd.read_sql(query, conn)
df.head()
```

	Id int64	EmployeeName o...	JobTitle object	BasePay float64	OvertimePay floa...	OtherPay float64	Benefits float64	TotalPay float64
0	36160	Gary Altenberg	Lieutenant, Fire Suppression	128808.87	220909.48	13126.31	44430.12	362844.6
1	36161	Gregory Suhr	Chief of Police	302578.0	0.0	18974.11	69810.19	321552.1
2	36162	Khoa Trinh	Electronic Maintenance Tech	111921.0	146415.32	78057.41	53102.29	336393.7

3	36163	Joanne Hayes-White	Chief, Fire Department	296943.01	0.0	17816.59	72047.88	314759.48
4	36164	Frederick Binkley	EMT/Paramedic/Firefighter	126863.19	192424.49	17917.18	44438.25	337204.81

Problem 3:

Further limit the table to the year 2012, employees making under 150,000, and sort in descending order by year.

```
query = 'SELECT * FROM salaries WHERE TotalPay < 150000 AND Year = 2012 ORDER BY Year DESC'

df = pd.read_sql(query, conn)
df.head()
```

	Id int64	EmployeeName object	JobTitle object	BasePay float64	OvertimePay float64	OtherPay float64	Benefits float64	TotalPay float64
1	37999	Tristan Levarado	Manager IV	145592.63	0.0	1500.0	61386.3	147092.63
0	37556	Marcia Bell	Law Librarian	140485.6	0.0	8387.1	69196.17	148872.7
3	38049	Gerardo Fries	Manager IV	145448.07	0.0	3486.0	58289.15	148934.07
2	38041	Paul Gambon	Manager V	145592.65	0.0	3486.0	58293.37	149078.65
4	38054	Masood Ordikhani	Manager IV	145606.89	0.0	3486.0	58007.96	149092.89

Aggregation

Problem 4:

Select the average base pay from the table.

```
query = 'SELECT AVG(BasePay) FROM Salaries'

df = pd.read_sql(query, conn)
df.head()
```

	AVG(BasePay) float64
0	66053.72928809702

Problem 5:

Produce and print the head of a dataframe that shows the average pay for each year (only use a single, simple query). Your result should have a column for the year and a column for the average base pay.

```
query = 'SELECT Year, AVG(BasePay) FROM Salaries GROUP BY Year'

df = pd.read_sql(query, conn)
df.head()
```

	Year int64	AVG(BasePay) float64
0	2011	63595.956516774524
1	2012	65436.40685742255
2	2013	68509.83215550765

Problem 6:

Create a dataframe with average base pay, benefits, and overtime for each job title, as well as a column with the total average.

```
query = 'SELECT AVG(BasePay),  AVG(Benefits), AVG(OvertimePay) FROM Salaries GROUP BY JobTitle'
df = pd.read_sql(query, conn)
df["TotalAveragePay"] = df['AVG(BasePay)'] + df['AVG(OvertimePay)'] + df['AVG(Benefits)']
df.head()
```

	AVG(BasePay) flo...	AVG(Benefits) flo...	AVG(OvertimePa...	TotalAveragePay f..	
0	43300.806506024106	0.0	373.200843373494	43674.0073493976	
1	46643.172	0.0	0.0	46643.172	
2	28732.66395833333	0.0	24.430625000000003	28757.0945833333	
3	62290.78	17975.59	0.0	80266.37	
4	66374.4	0.0	0.0	66374.4	

Table Creation

Problem 7:

Now we'll create our own table in our database. Separate the Salaries table by Year, and add it back to the database

```
for y in ['2011', '2012', '2013', '2014']:
    query = "SELECT * FROM salaries GROUP BY Year"

    df = pd.read_sql(query, conn)
    df.to_sql(name='Y'+y, con=conn, if_exists='replace')
```

Table Joining

Problem 8:

We'll move on to a new dataset for the next steps. Download the dataset from here (<https://www.kaggle.com/datasets/luizpaulodeoliveira/imdb-project-sql>) and load the sqlite file same as before. Start by just selecting everything in the movie to see what it looks like.

```
conn = sqlite3.connect("movies.sqlite")
query = 'SELECT * FROM movies'

df = pd.read_sql(query, conn)
df.head()
```

	id int64	original_title object	budget int64	popularity int64	release_date obj...	revenue int64	title object	vote_average flo...
0	43597	Avatar	237000000	150	2009-12-10	2787965087	Avatar	7.2
1	43598	Pirates of the Caribbean: At...	300000000	139	2007-05-19	961000000	Pirates of the Caribbean: At...	6.9
2	43599	Spectre	245000000	107	2015-10-26	880674609	Spectre	6.3
3	43600	The Dark Knight Rises	250000000	112	2012-07-16	1084939099	The Dark Knight Rises	7.6
4	43601	John Carter	260000000	43	2012-03-07	284139100	John Carter	6.1

Problem 9:

Create a dataframe that includes the entire contents of movies as well as the director's name.

```
query = 'SELECT movies.*, directors.name FROM movies JOIN directors ON movies. director_id = directors.id'
df = pd.read_sql(query, conn)

print(df.head(5))
df.size
```

3	2012-07-16	1084939099	The Dark Knight Rises
4	2012-03-07	284139100	John Carter

	vote_average	vote_count	\
0	7.2	11800	
1	6.9	4500	
2	6.3	4466	
3	7.6	9106	
4	6.1	2124	

	overview	\
0	In the 22nd century, a paraplegic Marine is di...	
1	Captain Barbossa, long believed to be dead, ha...	
2	A cryptic message from Bond's past sends him o...	
3	Following the death of District Attorney Harve...	
4	John Carter is a war-weary, former military ca...	

	tagline	uid	director_id	\
0	Enter the World of Pandora.	19995	4762	
1	At the end of the world, the adventure begins.	285	4763	
2	A Plan No One Escapes	206647	4764	
3	The Legend Ends	49026	4765	
4	Lost in our world, found in another.	49529	4766	

	name
0	James Cameron
1	Gore Verbinski
2	Sam Mendes
3	Christopher Nolan
4	Andrew Stanton

66822

Analysis

The next few problems will be more involved! You'll need to combine some concepts you've learned. For each cell, show your work.

Problem 10:

What is the average budget used for the top 10 grossing movies?

```
query = 'SELECT budget FROM movies ORDER by revenue DESC LIMIT 10;'

df = pd.read_sql(query, conn)

df['budget'].mean()
```

195100000.0

Problem 11:

Which directors have the highest voting average? - show the top 5 directors' name and their average rating

```
query = '''
    SELECT directors.name, movies.vote_average AS avg_rating
    FROM movies
    JOIN directors ON movies.director_id = directors.id
    GROUP BY movies.director_id
    ORDER BY avg_rating DESC
    LIMIT 5
    '''

df = pd.read_sql(query, conn)

df
```

	name object	avg_rating float64	
0	Gary Sinyor	10.0	
1	Rohit Jugraj	9.5	
2	Lance Hool	9.3	
3	Floyd Mutrux	8.5	
4	John Cromwell	8.4	

Problem 12:

What are the top five directors by average budget?

```
query = '''
    SELECT AVG(movies.budget) AS avg_budget, directors.name
    FROM movies
    JOIN directors ON movies.director_id = directors.id
    GROUP BY directors.name
    ORDER BY AVG(movies.budget) DESC

    LIMIT 5 '''

df = pd.read_sql(query, conn)

print(df.head(5))
```

	avg_budget	name
0	2.600000e+08	Byron Howard
1	2.000000e+08	Dan Scanlon
2	2.000000e+08	Lee Unkrich
3	1.933333e+08	David Yates
4	1.850000e+08	Brenda Chapman